

MuscleHub A/B Test

Step 1: Get started with SQL

Like most businesses, Janet keeps her data in a SQL database. Normally, you'd download the data from her database to a csv file, and then load it into a Jupyter Notebook using Pandas.

For this project, you'll have to access SQL in a slightly different way. You'll be using a special Codecademy library that lets you type SQL queries directly into this Jupyter notebook. You'll have to pass each SQL query as an argument to a function called `sql_query`. Each query will return a Pandas DataFrame.

```
In [1]: from codecademySQL import sql_query
```

Step 2: Get your dataset

Let's get started!

Janet of MuscleHub has a SQLite database, which contains several tables that will be helpful to you in this investigation:

- `visits` contains information about potential gym customers who have visited MuscleHub
- `fitness_tests` contains information about potential customers in "Group A", who were given a fitness test
- `applications` contains information about any potential customers (both "Group A" and "Group B") who filled out an application. Not everyone in `visits` will have filled out an application.
- `purchases` contains information about customers who purchased a membership to MuscleHub.

```
In [2]: #examine visits
sql_query('''
SELECT *
FROM visits
LIMIT 5
''')
```

```
Out[2]:   index first_name last_name          email    gender visit_date
0        0      Karen  Manning  Karen.Manning@gmail.com  female   5-1-17
1        1     Annette    Boone  AB9982@gmail.com  female   5-1-17
2        2   Salvador   Merritt  SalvadorMerritt12@outlook.com    male   5-1-17
3        3     Martha  Maxwell  Martha.Maxwell@gmail.com  female   5-1-17
4        4      Andre   Mayer  AndreMayer90@gmail.com    male   5-1-17
```

```
In [3]: #examine fitness_tests
sql_query('''
SELECT *
FROM fitness_tests
LIMIT 5
''')
```

```
Out[3]:   index first_name last_name          email gender fitness_test_date
0         0      Kim    Walter KimWalter58@gmail.com female 2017-07-03
1         1      Tom  Webster  TW3857@gmail.com male 2017-07-02
2         2  Marcus     Bauer Marcus.Bauer@gmail.com male 2017-07-01
3         3  Roberta     Best  RB6305@hotmail.com female 2017-07-02
4         4   Carrie    Francis CF1896@hotmail.com female 2017-07-05
```

```
In [4]: #examine applications
sql_query('''
SELECT *
FROM applications
LIMIT 5
''')
```

```
Out[4]:   index first_name last_name          email gender application_date
0         0      Roy    Abbott RoyAbbott32@gmail.com male 2017-08-12
1         1    Agnes Acevedo AgnesAcevedo1@gmail.com female 2017-09-29
2         2  Roberta Acevedo  RA8063@gmail.com female 2017-09-15
3         3    Darren   Acosta DAcosta1996@hotmail.com male 2017-07-26
4         4    Vernon   Acosta VAcosta1975@gmail.com male 2017-07-14
```

```
In [5]: #examine purchases
sql_query('''
SELECT *
FROM purchases
LIMIT 5
''')
```

```
Out[5]:   index first_name last_name          email gender purchase_date
0         0      Roy    Abbott RoyAbbott32@gmail.com male 2017-08-18
1         1  Roberta Acevedo  RA8063@gmail.com female 2017-09-16
2         2    Vernon   Acosta VAcosta1975@gmail.com male 2017-07-20
3         3    Darren   Acosta DAcosta1996@hotmail.com male 2017-07-27
4         4      Dawn  Adkins Dawn.Adkins@gmail.com female 2017-08-24
```

We'd like to download a giant DataFrame containing all of this data. You'll need to write a query

that does the following things:

1. Not all visits in `visits` occurred during the A/B test. You'll only want to pull data where `visit_date` is on or after `7-1-17`.
2. You'll want to perform a series of `LEFT JOIN` commands to combine the four tables that we care about. You'll need to perform the joins on `first_name`, `last_name`, and `email`. Pull the following columns:
 - `visits.first_name`
 - `visits.last_name`
 - `visits.gender`
 - `visits.email`
 - `visits.visit_date`
 - `fitness_tests.fitness_test_date`
 - `applications.application_date`
 - `purchases.purchase_date`

Save the result of this query to a variable called `df`.

Hint: your result should have 5004 rows. Does it?

```
In [6]: df = sql_query('''
SELECT visits.first_name,
       visits.last_name,
       visits.gender,
       visits.email,
       visits.visit_date,
       fitness_tests.fitness_test_date,
       applications.application_date,
       purchases.purchase_date
FROM visits
LEFT JOIN fitness_tests
    ON visits.first_name = fitness_tests.first_name
    AND visits.last_name = fitness_tests.last_name
    AND visits.email = fitness_tests.email
LEFT JOIN applications
    ON visits.first_name = applications.first_name
    AND visits.last_name = applications.last_name
    AND visits.email = applications.email
LEFT JOIN purchases
    ON visits.first_name = purchases.first_name
    AND visits.last_name = purchases.last_name
    AND visits.email = purchases.email
WHERE visits.visit_date >= '7-1-17'
''')
df.tail()
#dataframe has 5004 total rows
```

	first_name	last_name	gender	email	visit_date	fitness_test_date	application
4999	Rachel	Hensley	female	RachelHensley38@gmail.com	9-9-17		None
5000	Leon	Harmon	male	Leon.Harmon@gmail.com	9-9-17	2017-09-15	

	first_name	last_name	gender		email	visit_date	fitness_test_date	application_date
5001	Andy	Pratt	male		AndyPratt27@gmail.com	9-9-17	2017-09-15	
5002	Ruben	Nielsen	male	RubenNielsen93@hotmail.com		9-9-17	None	2017-09-15
5003	Charles	Carver	male		CC2490@gmail.com	9-9-17	2017-09-12	

Step 3: Investigate the A and B groups

We have some data to work with! Import the following modules so that we can start doing analysis:

- import pandas as pd
- from matplotlib import pyplot as plt

In [7]:

```
import pandas as pd
from matplotlib import pyplot as plt
```

We're going to add some columns to df to help us with our analysis.

Start by adding a column called ab_test_group . It should be A if fitness_test_date is not None , and B if fitness_test_date is None .

In [8]:

```
df["ab_test_group"] = df.fitness_test_date.apply(lambda x: "A" if pd.notnull(x) else "B")
df.head()
```

Out[8]:

	first_name	last_name	gender		email	visit_date	fitness_test_date	application_date
0	Kim	Walter	female		KimWalter58@gmail.com	7-1-17	2017-07-03	None
1	Tom	Webster	male		TW3857@gmail.com	7-1-17	2017-07-02	None
2	Edward	Bowen	male	Edward.Bowen@gmail.com		7-1-17	None	2017-07-02
3	Marcus	Bauer	male	Marcus.Bauer@gmail.com		7-1-17	2017-07-01	2017-07-01
4	Roberta	Best	female		RB6305@hotmail.com	7-1-17	2017-07-02	None

Let's do a quick sanity check that Janet split her visitors such that about half are in A and half are in B.

Start by using groupby to count how many users are in each ab_test_group . Save the results to ab_counts .

In [9]:

```
ab_counts = df.groupby("ab_test_group").first_name.count().reset_index()
ab_counts
```

Out[9]:

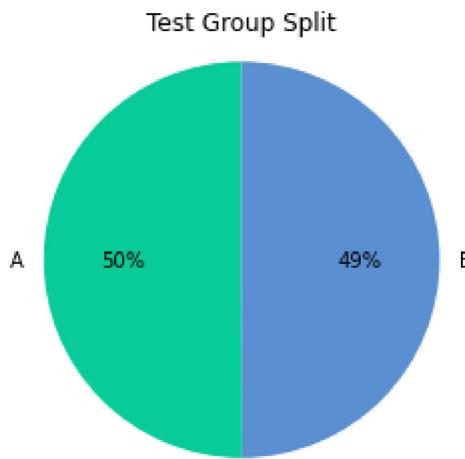
	ab_test_group	first_name
0	A	2504

ab_test_group	first_name
1	B 2500

We'll want to include this information in our presentation. Let's create a pie chart using `plt.pie`. Make sure to include:

- Use `plt.axis('equal')` so that your pie chart looks nice
- Add a legend labeling A and B
- Use `autopct` to label the percentage of each group
- Save your figure as `ab_test_pie_chart.png`

```
In [10]: plt.pie(ab_counts.first_name.values, labels=["A", "B"], autopct="%d%%", startangle=90, plt.axis("equal") plt.title("Test Group Split") plt.show() plt.savefig("ab_test_pie_chart.png")
```



<Figure size 432x288 with 0 Axes>

Step 4: Who picks up an application?

Recall that the sign-up process for MuscleHub has several steps:

1. Take a fitness test with a personal trainer (only Group A)
2. Fill out an application for the gym
3. Send in their payment for their first month's membership

Let's examine how many people make it to Step 2, filling out an application.

Start by creating a new column in `df` called `is_application` which is `Application` if `application_date` is not `None` and `No Application`, otherwise.

```
In [11]: df["is_application"] = df.application_date.apply(lambda x: "Application" if pd.notnull(
```

Now, using `groupby`, count how many people from Group A and Group B either do or don't pick up an application. You'll want to group by `ab_test_group` and `is_application`. Save this new

DataFrame as `app_counts`

```
In [12]: app_counts = df.groupby(["ab_test_group", "is_application"]).first_name.count().reset_index()
app_counts
```

```
Out[12]: ab_test_group  is_application  first_name
0               A      Application       250
1               A    No Application     2254
2               B      Application       325
3               B    No Application     2175
```

We're going to want to calculate the percent of people in each group who complete an application. It's going to be much easier to do this if we pivot `app_counts` such that:

- The index is `ab_test_group`
- The columns are `is_application`. Perform this pivot and save it to the variable `app_pivot`. Remember to call `reset_index()` at the end of the pivot!

```
In [13]: app_pivot = app_counts.pivot(index="ab_test_group", columns="is_application", values="first_name")
```

Define a new column called `Total`, which is the sum of `Application` and `No Application`.

```
In [14]: app_pivot["Total"] = app_pivot["Application"] + app_pivot["No Application"]
```

Calculate another column called `Percent with Application`, which is equal to `Application` divided by `Total`.

```
In [15]: app_pivot["Percent with Application"] = app_pivot["Application"] / app_pivot["Total"]
app_pivot
```

```
Out[15]: is_application  ab_test_group  Application  No Application  Total  Percent with Application
0               A               250           2254      2504        0.09984
1               B               325           2175      2500        0.13000
```

It looks like more people from Group B turned in an application. Why might that be?

We need to know if this difference is statistically significant.

Choose a hypothesis tests, import it from `scipy` and perform it. Be sure to note the p-value. Is this result significant?

```
In [16]: from scipy.stats import chi2_contingency

app_contingency = ([250, 2254], [325, 2175])
_, app_p, _, _ = chi2_contingency(app_contingency)
if app_p < .05:
```

```

        print(round(app_p, 5))
        print("Significant!")
else:
    print(round(app_p, 5))
    print("Not Significant!")

```

0.00096
Significant!

Step 5: Who purchases a membership?

Of those who picked up an application, how many purchased a membership?

Let's begin by adding a column to `df` called `is_member` which is Member if `purchase_date` is not `None`, and Not Member otherwise.

In [17]:

```
df["is_member"] = df.purchase_date.apply(lambda x: "Member" if pd.notnull(x) else "Not Member")
```

Now, let's create a DataFrame called `just_apps` the contains only people who picked up an application.

In [18]:

```
just_apps = df[df.is_application == "Application"]
just_apps.head()
```

Out[18]:

	first_name	last_name	gender		email	visit_date	fitness_test_date	application_date
2	Edward	Bowen	male	Edward.Bowen@gmail.com	7-1-17	None	2017-07-01	2017-07-01
3	Marcus	Bauer	male	Marcus.Bauer@gmail.com	7-1-17	2017-07-01	2017-07-01	2017-07-01
9	Salvador	Cardenas	male	SCardenas1980@gmail.com	7-1-17	2017-07-07	2017-07-07	2017-07-07
11	Valerie	Munoz	female	VMunoz1998@gmail.com	7-1-17	2017-07-03	2017-07-03	2017-07-03
35	Michael	Burks	male	MB9820@gmail.com	7-1-17	None	2017-07-01	2017-07-01

Great! Now, let's do a `groupby` to find out how many people in `just_apps` are members and aren't members from each group. Follow the same process that we did in Step 4, including pivoting the data. You should end up with a DataFrame that looks like this:

is_member	ab_test_group	Member	Not Member	Total	Percent Purchase
0	A	?	?	?	?
1	B	?	?	?	?

Save your final DataFrame as `member_pivot`.

In [19]:

```
member_counts = just_apps.groupby(["ab_test_group", "is_member"]).first_name.count().reset_index()
#member_count

member_pivot = member_counts.pivot(index="ab_test_group", columns="is_member", values="count")
```

```

member_pivot["Total"] = member_pivot["Member"] + member_pivot["Not Member"]
member_pivot["Percent Purchased"] = member_pivot["Member"] / member_pivot["Total"]
member_pivot

```

Out[19]:

is_member	ab_test_group	Member	Not Member	Total	Percent Purchased
0	A	200	50	250	0.800000
1	B	250	75	325	0.769231

It looks like people who took the fitness test were more likely to purchase a membership **if** they picked up an application. Why might that be?

Just like before, we need to know if this difference is statistically significant. Choose a hypothesis tests, import it from `scipy` and perform it. Be sure to note the p-value. Is this result significant?

In [20]:

```

member_contingency = ([200, 50], [250, 75])
_, mem_p, _, _ = chi2_contingency(member_contingency)
if mem_p < .05:
    print(round(mem_p, 5))
    print("Significant!")
else:
    print(round(mem_p, 5))
    print("Not Significant!")

```

0.43259
Not Significant!

Previously, we looked at what percent of people **who picked up applications** purchased memberships. What we really care about is what percentage of **all visitors** purchased memberships. Return to `df` and do a `groupby` to find out how many people in `df` are members and aren't members from each group. Follow the same process that we did in Step 4, including pivoting the data. You should end up with a DataFrame that looks like this:

is_member	ab_test_group	Member	Not Member	Total	Percent Purchase
0	A	?	?	?	?
1	B	?	?	?	?

Save your final DataFrame as `final_member_pivot`.

In [21]:

```

final_member_count = df.groupby(["ab_test_group", "is_member"]).first_name.count().reset_index()
final_member_pivot = final_member_count.pivot(index="ab_test_group", columns="is_member")
final_member_pivot["Total"] = final_member_pivot["Member"] + final_member_pivot["Not Member"]
final_member_pivot["Percent Purchased"] = final_member_pivot["Member"] / final_member_pivot["Total"]
final_member_pivot

```

Out[21]:

is_member	ab_test_group	Member	Not Member	Total	Percent Purchased
0	A	200	2304	2504	0.079872
1	B	250	2250	2500	0.100000

Previously, when we only considered people who had **already picked up an application**, we saw that there was no significant difference in membership between Group A and Group B.

Now, when we consider all people who **visit MuscleHub**, we see that there might be a significant different in memberships between Group A and Group B. Perform a significance test and check.

In [22]:

```
fmem_contingency = ([200, 2304], [250, 2250])
_, fmem_p, _, _ = chi2_contingency(fmem_contingency)
if fmem_p < .05:
    print(round(fmem_p, 5))
    print("Significant!")
else:
    print(round(fmem_p, 5))
    print("Not Significant!")
```

```
0.01472
Significant!
```

Step 6: Summarize the acquisition funnel with a chart

We'd like to make a bar chart for Janet that shows the difference between Group A (people who were given the fitness test) and Group B (people who were not given the fitness test) at each state of the process:

- Percent of visitors who apply
- Percent of applicants who purchase a membership
- Percent of visitors who purchase a membership

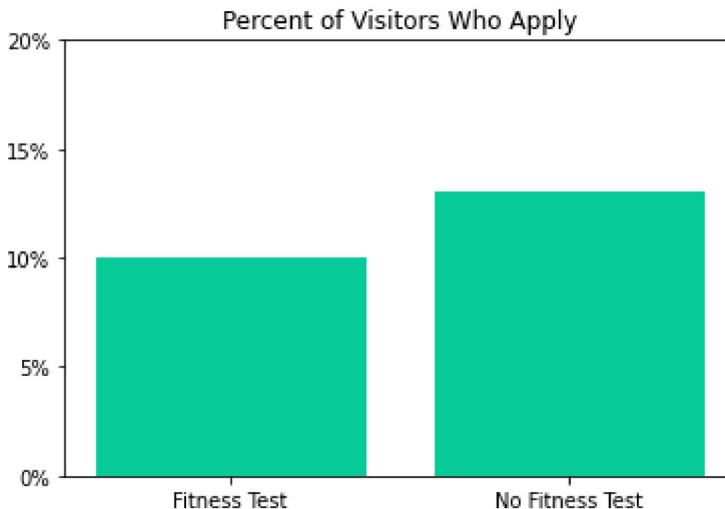
Create one plot for **each** of the three sets of percentages that you calculated in `app_pivot` , `member_pivot` and `final_member_pivot` . Each plot should:

- Label the two bars as `Fitness Test` and `No Fitness Test`
- Make sure that the y-axis ticks are expressed as percents (i.e., 5%)
- Have a title

In [23]:

```
#Percent of visitors who apply
ax = plt.subplot()
plt.bar(range(len(app_pivot)), app_pivot["Percent with Application"].values, color="#075690")
ax.set_xticks(range(len(app_pivot)))
ax.set_xticklabels(["Fitness Test", "No Fitness Test"])
ax.set_yticks([0, 0.05, 0.10, 0.15, 0.20])
ax.set_yticklabels(["0%", "5%", "10%", "15%", "20%"])
plt.title("Percent of Visitors Who Apply")
plt.show()
plt.savefig("pct_visitors_who_apply.png")

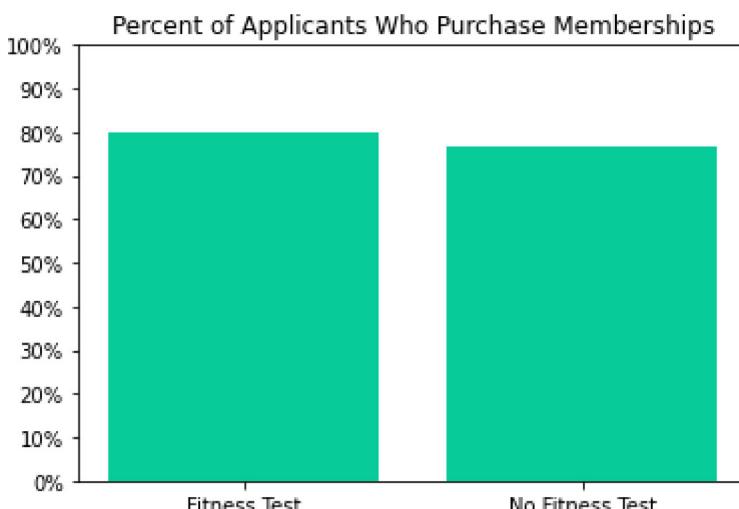
app_sig = "The difference of more applicants from group b is significant given p-value:
app_sig
```



```
Out[23]: 'The difference of more applicants from group b is significant given p-value: 0.00096'
<Figure size 432x288 with 0 Axes>
```

```
In [24]: #Percent of applicants who purchase a membership
ax = plt.subplot()
plt.bar(range(len(member_pivot)), member_pivot["Percent Purchased"].values, color="#07C")
ax.set_xticks(range(len(member_pivot)))
ax.set_xticklabels(["Fitness Test", "No Fitness Test"])
ax.set_yticks([0, 0.1, 0.2, 0.3, 0.4, 0.5, 0.6, 0.7, 0.8, 0.9, 1])
ax.set_yticklabels(["0%", "10%", "20%", "30%", "40%", "50%", "60%", "70%", "80%", "90%", "100%"])
plt.title("Percent of Applicants Who Purchase Memberships")
plt.show()
plt.savefig("pct_applicants_who_purchase_memberships.png")

mem_sig = "The difference of more applicants from group a purchasing memberships is not significant"
mem_sig
```

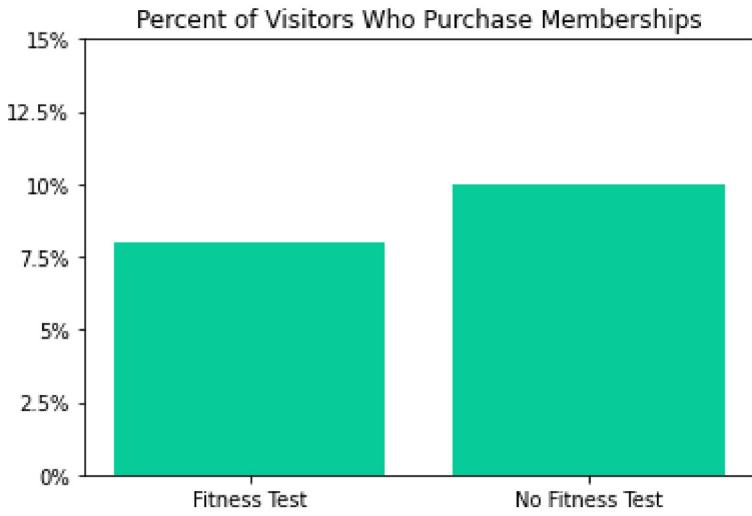


```
Out[24]: 'The difference of more applicants from group a purchasing memberships is not significant given p-value: 0.4326'
<Figure size 432x288 with 0 Axes>
```

```
In [25]: #Percent of visitors who purchase a membership
ax = plt.subplot()
plt.bar(range(len(final_member_pivot)), final_member_pivot["Percent Purchased"].values,
ax.set_xticks(range(len(final_member_pivot)))
```

```
ax.set_xticklabels(["Fitness Test", "No Fitness Test"])
ax.set_yticks([0, 0.025, 0.05, 0.075, 0.1, 0.125, 0.15])
ax.set_yticklabels(["0%", "2.5%", "5%", "7.5%", "10%", "12.5%", "15%"])
plt.title("Percent of Visitors Who Purchase Memberships")
plt.show()
plt.savefig("pct_visitors_who_purchase_memberships.png")

fmem_sig = "The difference of more memberships purchased by group b is significant p-value: 0.0147"
```



```
Out[25]: 'The difference of more memberships purchased by group b is significant p-value: 0.0147'
<Figure size 432x288 with 0 Axes>
```